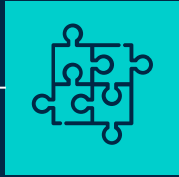


Quick! Draw! Doodle Recognition

By Marcus Tan

TABLE OF CONTENTS



01

Problem Statement
& EDA



02

CNN model &
Transfer
learning



03

Conclusion
&
Recommendations

Problem Statement & EDA

01


Problem Statement

- Dataset of over 50 million hand drawn drawings with 340 labels
- Tasked by Google to build a better classifier
- Challenge is to predict the drawings within 3 guesses
- Metric - Accuracy & Top 3 accuracy
- Immediate impact on handwriting recognition applications (e.g. Automatic Speech Recognition)

Datasets




- Train set (340 labels - totalling 50 million+ drawings)
- Test set (340 labels - totalling 100k+ drawings)
- Dataset has been preprocessed by kaggle
- No cleaning needed, No missing values
- Dataset contains drawing strokes, labels and recognized (True/False)

Drawing examples






- Left is recognized as False
- Right is recognized as True
- Left one is oddly shaped and random strokes

More drawing examples (True)



- All drawings are recognized as True
- Red arrows show different direction of drawing
- Form of Data Augmentation
- Helpful in generalizing to new images and reduce variance

More drawing examples (False)




- All drawings are recognized as False
- Most of the drawings do not fit the label
- Exception - bird on the left
- Not including False drawings into the model
- May confuse the model

CNN model & Transfer learning

02

Basic CNN Model Architecture



Basic CNN Model Architecture

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	320
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_1 (Conv2D)	(None, 16, 16, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
flatten (Flatten)	(None, 4096)	0
dense (Dense)	(None, 680)	2785960
dropout (Dropout)	(None, 680)	0
dense_1 (Dense)	(None, 340)	231540

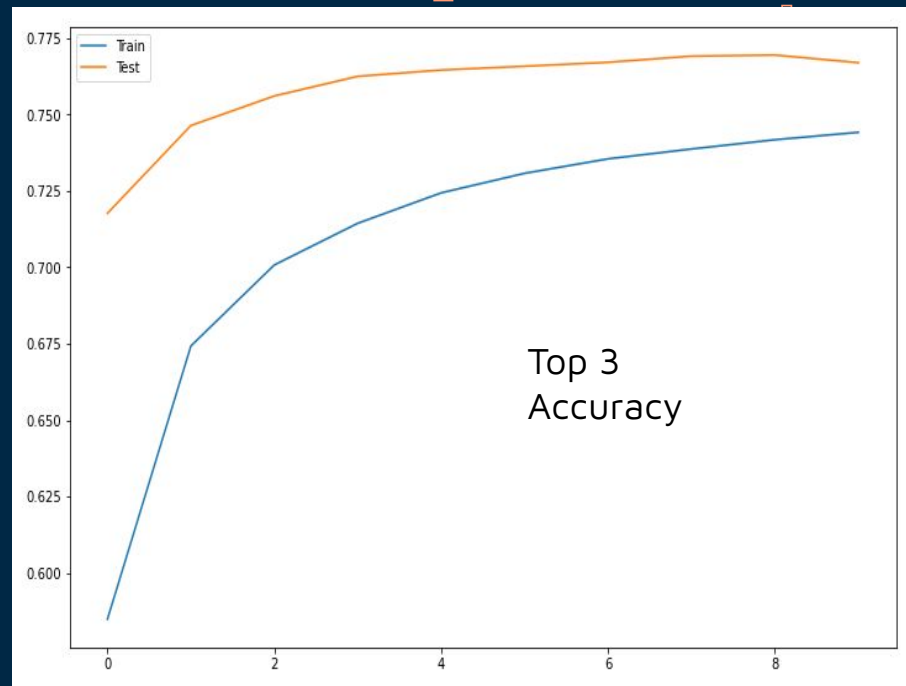

=====
Total params: 3,036,316

Trainable params: 3,036,316

Non-trainable params: 0

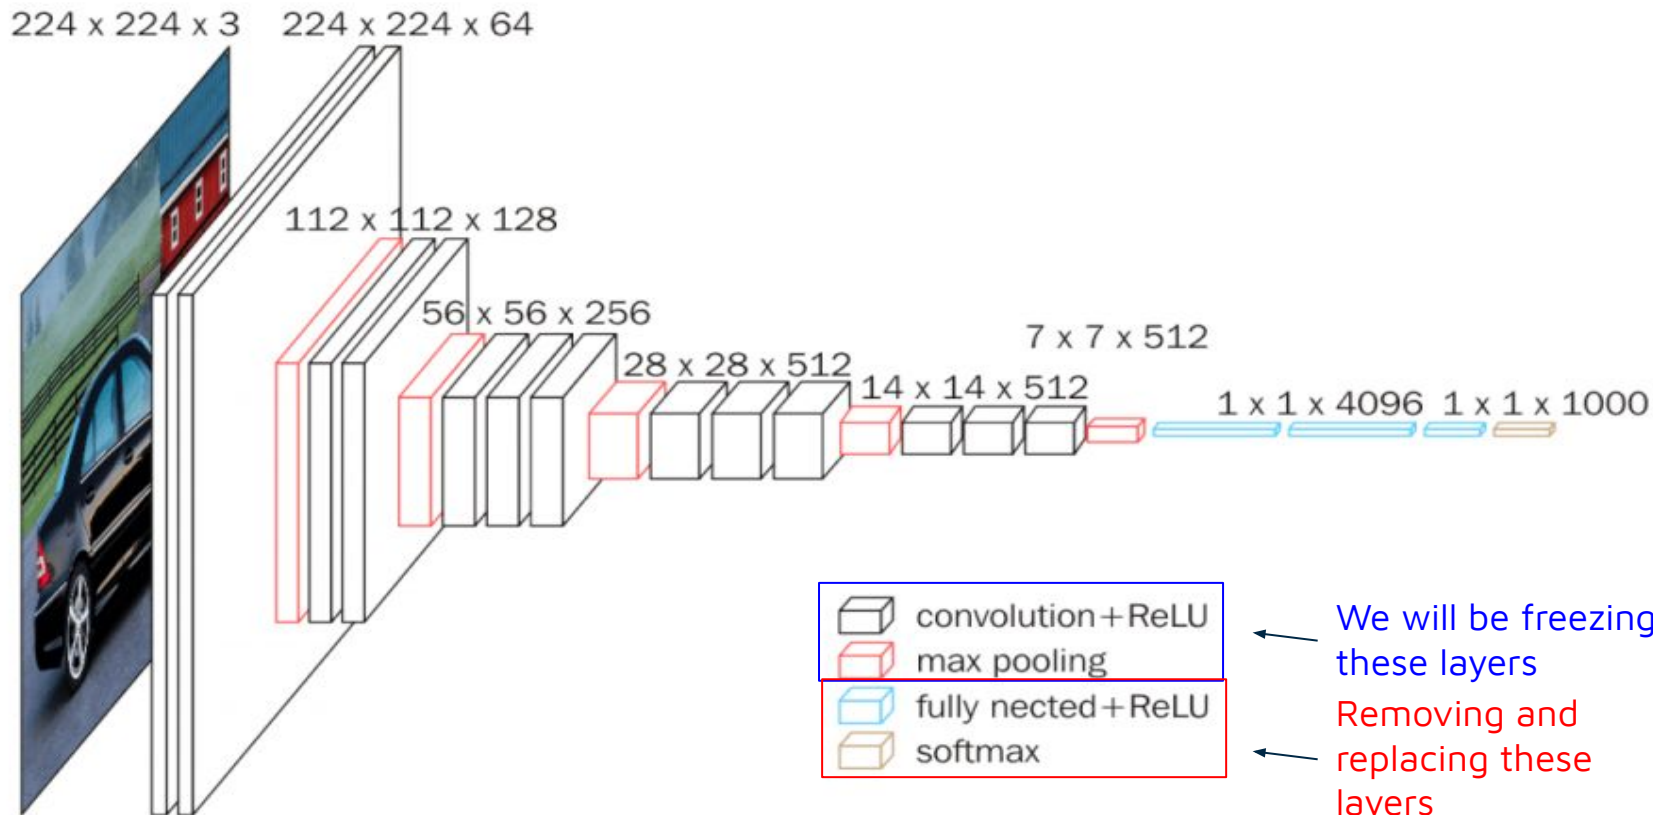
- Image fed in as (32, 32, 1)
- Box in red - giving the 3D image depth and extracting features
- Box in blue - flattens into a 1 dimensional array and gives an output of 340 probabilities
- Not too many layers as training time increases exponentially
- Used 2000 image per class due to limited RAM

Basic CNN Model Performance



- Validation Accuracy of 0.56, Top 3 Accuracy of 0.76
- No overfitting - lines slowly converging
- Test is performing better - low variance

Transfer Learning - VGG16



Transfer Learning - VGG16

Model: "sequential"

Layer (type)	Output Shape	Param #
3layers_for_vgg (Conv2D)	(None, 32, 32, 3)	30
vgg16 (Functional)	(None, 1, 1, 512)	14714688
flatten (Flatten)	(None, 512)	0
dense (Dense)	(None, 680)	348840
dropout (Dropout)	(None, 680)	0
dense_1 (Dense)	(None, 340)	231540

Total params: 15,295,098
Trainable params: 580,410
Non-trainable params: 14,714,688

- Image fed in as (32, 32, 1)
- Box in red - Created a Conv2D layer before VGG16 base model to create 'RGB' image
- Box in blue - flattens into a 1 dimensional array and gives an output of 340 probabilities
- Trained the model on Google Colab
- Used 1000 image per class due to limited RAM


Summary of Model score

Models	Number of image used per class	Train Top 3 Accuracy	Test Top 3 Accuracy	Kaggle Top 3 Accuracy
Basic CNN	2000	0.74	0.76	0.61
VGG16 (32x32)	1000	0.55	0.57	Did not try
VGG16 (128x128)	75 x 5	0.76	0.78	0.62

- Lesser image used as dimension gets bigger due to limited RAM
- Bigger dimension better score for pre-trained models (VGG16 trained on 224x224x3 images)
- Kaggle score did not perform well - likely due to model not trained on enough images
- Trained on less than 1 million images out of 50 million+
- VGG16 (128x128) model saved, loaded and trained on 5 different sets of 75 images per class

Challenges faced during project

- Basic model run on local CPU took more than 12hrs
- Google colab RAM limit - reduce images in read in
- Pre-trained model overfits very quickly
- Reduce start Learning Rate by a factor of 0.25 to prevent overfitting
- Lastly a lot of rabbit hole



Conclusion & Recommendations

03

Conclusion and Recommendations

- Best model - VGG16 (128x128)
- Model predicts relatively well with 0.62 top 3 accuracy
- Limitations:
 - Only able to predict images from the 340 labels
 - Will predict colored drawings as if it were grayscale
- Further improvements:
 - More time and computing power
 - Read in images as 224x224
 - Train on all 50 million+ images
 - Try out other pre-trained models (MobileNet, EfficientNet)
- Recommendations:
 - Toddler programme/application to guide them in drawing right
 - Learn how to draw and develop artistic talents at an early age

Do you have any questions?



THANK YOU



CREDITS: This presentation template was created by [Slidesgo](#),
including icons by [Flaticon](#), and infographics & images by [Freepik](#)
Please keep this slide for attribution